Cloud cost optimization typically follows two main approaches. The first is usage optimization, also known as workload optimization, which focuses on reducing the amount of cloud resources consumed while still maintaining the performance. This includes practices like rightsizing instances, scheduling workloads, and aligning consumption with demand patterns. The other is rate optimization, which reduces the unit cost of cloud resources by leveraging discount instruments such as Reserved Instances and Savings Plans.

Both approaches are necessary, but rate optimization requires a deeper understanding before any commitments are made. Choosing the wrong plan, overcommitting, or letting discounts expire without renewal can lead to inefficiencies and missed savings.

This blog will help you cover the basics of AWS Savings Plans, how they work, the key differences between different SPs offered, and how to make the most of them without compromising agility. (If you're new to rate-based savings strategies, start with our 101 guide to AWS Reserved Instances.)

What Are AWS Savings Plans?

AWS Savings Plans are long-term pricing models that offer up to 72% cost savings compared to on-demand rates in exchange for a one- or three-year commitment to a consistent level of usage, measured in dollars per hour. 72% represents the highest possible discount and applies only under specific conditions such as EC2 Instance Savings Plans with a three-year All Upfront payment. Actual savings will vary depending on the plan type, term, payment option, and covered services.

For Savings Plans, certain configurations must be defined at purchase — such as plan type, term length, payment option, and in the case of EC2 Instance Savings Plans, the specific instance family and region.

The big shift with Savings Plans is the unit of commitment. With RIs, you make instance-specific commitments (e.g., quantity 1, c8g.xlarge running Linux with shared tenancy in the us-west-2 region for one year), which can be challenging to plan for and manage. With Savings Plans, you instead make an aggregate hourly dollar commitment (e.g. $9.57/hr for one or three years), and then AWS applies discounts automatically on eligible usage.

Managing Savings Plan commitments

Additionally, AWS now offers the ability to schedule Savings Plan purchases. They can be queued or scheduled to occur on a future date, down to the exact second for coverage to begin.

Even though Savings Plans are based on a fixed dollar-per-hour commitment, selecting the right configuration is still critical. You'll need to define parameters like plan type, term length, and payment option, all of which impact how coverage is applied and how much flexibility you retain. If your usage exceeds the commitment amount, the excess is charged at On-Demand rates. If your usage falls below the commitment, the unused portion goes to waste, making right-sizing and ongoing monitoring key to maximizing investment value.

Benefits of AWS Savings Plans

For organizations committing to sustained AWS usage, Savings Plans offer long-term savings with less operational complexity. Below are the key benefits they bring to cost and resource management:

Meaningful cost savings

Savings Plans offer up to 72% in savings compared to On-Demand pricing. They provide a more horizontally flexible way to reduce cloud spend without compromising scalability or performance. Horizontal flexibility refers to the ability of Compute Savings Plans to cover usage across services, regions, and workloads.

Automatic application across services

Once purchased, Savings Plans automatically apply to eligible usage without requiring manual allocation. This reduces management overhead and ensures discounts are continuously applied where possible.

Spend-based flexibility

Savings Plans are based on a fixed hourly spend, meaning the discount continues to apply even as instance types, regions, or operating systems change (with the exception of EC2 Instance Savings Plans). This flexibility makes it easier to manage commitments as infrastructure evolves, especially for teams in the early stages of their FinOps maturity.

On-Demand vs. AWS Savings Plans

On-Demand pricing offers maximum flexibility, allowing you to pay per second with no long-term commitment. This makes it a good fit for unpredictable or short-lived workloads.

Savings Plans, in contrast, require a one- or three-year commitment to a consistent hourly spend, but in return, they offer significantly lower rates. In optimal scenarios, discounts can reach up to 72%, though actual savings depend on factors such as plan type, payment method, term length, and the services covered. Savings Plans are better suited for stable, long-running workloads where usage patterns are more predictable.

When To Use AWS Savings Plans

Use AWS Savings Plans when you can commit to a consistent level of AWS spend over time but want the flexibility to change the workloads to which the savings apply. They’re ideal for environments with steady usage patterns but evolving architectures, such as organizations migrating from EC2 to Fargate or shifting between instance families.

Compared to Standard Reserved Instances, which are better suited for fixed, known configurations, Savings Plans work best when you want long-term savings without locking into specific instance types or regions. They also simplify cost management for multi-service compute strategies and are often easier to operationalize at scale.

How Do AWS Savings Plans Work?

AWS Savings Plans apply automatically to eligible usage, offering discounted rates in exchange for a commitment to a consistent hourly spend over a one- or three-year term. The hourly commitment amount must be specified in post-discount dollars.

In other words, not the amount of on-demand coverage you want, but the amount you'll pay after the SP discount is applied. This can be difficult to predict and calculate, but there are tools that can assist.

Here's how Savings Plans are purchased and applied:

> Purchase options: Savings Plans are available for purchase in both the AWS Console and via API.
> Automatic application: Once purchased, the plan tracks your hourly usage across supported services and automatically applies the discount where it yields the greatest savings first.
> Prioritization: This means usage in regions or configurations with higher discount potential is prioritized.
> Beyond commitment: Any usage beyond that is billed at standard On-Demand rates.

Although Savings Plans provide great flexibility within the specified dollar-per-hour commitment, certain configurations must be defined at purchase, such as plan type (Compute, EC2 Instance, or SageMaker), term length, payment option, and in the case of EC2 Instance Savings Plans, the

specific instance family and region. These parameters cannot be changed once committed.

The value of a Savings Plan depends heavily on how closely your actual usage aligns with the committed spend.

> Overcommitment leads to wasted spend, as any unused portion of the hourly commitment is lost.
> Undercommitment means you'll miss out on available discounts, paying On-Demand rates for usage that could have been covered.

What Are the Payment Options for AWS Savings Plans?

To enhance financial flexibility, AWS Savings Plans offer three different payment options to choose from:

> All Upfront: Pay for your entire commitment at the beginning and receive the largest discount possible.
> Partial Upfront: Pay a partial portion of your commitment at the beginning and the rest on a monthly basis for moderate discounts.
> No Upfront: Pay nothing upfront. Instead, make payments on your commitment in arrears, on a monthly basis. This option requires no cash upfront, but offers the lowest discounts.

Types of AWS Savings Plans

AWS offers three types of Savings Plans: Compute Savings Plans, EC2 Instance Savings Plans and Sagemaker Savings Plans. Each is designed to balance savings and flexibility differently, depending on how predictable your workloads are and how much infrastructure change you anticipate over time. Here's a detailed overview of the three:

EC2 Instance Savings Plans

EC2 Instance Savings Plans are a type of AWS Savings Plan that offer up to 72% discount compared to On-Demand pricing. These plans apply to EC2 usage within a specific instance family and region, while allowing flexibility across instance sizes, operating systems, and tenancies. Once purchased, the discount automatically applies to all matching EC2 usage up to the committed hourly spend.

For example, a $1.00/hr commitment for m7i in us-east-2 can automatically cover different combinations, such as multiple m7i.2xlarge Linux instances or a single m7i.4xlarge Windows SQL Server instance. For those familiar with instance-size flexibility (ISF), you can think of the EC2 Instance Savings Plan like “super ISF.” Since Savings Plans are dollar- vs. instance-based, they can float across operating systems and tenancy in addition to instance size.

Key characteristics of EC2 Instance Savings Plans:

Instance-family and region scope: Applies only to EC2 usage within the chosen instance family and region(e.g., m8g in N.Virginia).
Limited flexibility: You cannot switch instance families or regions and retain the discount, but you can move between instance sizes, operating systems, and tenancies within the selected family.
Higher savings: Offers better discounts than Compute Savings Plans due to the narrower scope of flexibility.
Automatic application: Within the defined parameters, the AWS billing system automatically applies discounts to matching EC2 usage.
ISPs apply to Commercial and GovCloud, plus Wavelength Zones.
Term and payment options: Available in one- or three-year terms with the same three payment model choices.
Best for: Steady-state EC2 workloads where the instance family and region are unlikely to change, such as legacy systems or workloads with predictable usage.

Compute Savings Plans

Compute Savings Plans are a type of AWS Savings Plan that offer up to 66% discount compared to On-Demand pricing. They provide broader flexibility than EC2 Instance Savings Plans, applying to any eligible compute usage (EC2, Fargate, Lambda) regardless of instance family, size, operating system, tenancy, or AWS region. This makes them well-suited for teams with dynamic workloads or evolving infrastructure strategies.

For example, a CSP commitment can automatically apply across EC2 instances or dedicated hosts running in different regions, across varying instance families, or even to containerized workloads running on AWS Fargate or functions executed via AWS Lambda, depending on which workload has the most discount potential at a given time.

One of the key advantages of Compute Savings Plans is their support for platform flexibility and global usage. As more teams shift to containerized and serverless architectures, these plans allow you to secure discounts today, even if you plan to move from EC2 to AWS Fargate or Lambda later.

The plan also spans all AWS regions, Commercial and GovCloud partitions, plus Local Zones and Wavelength Zones, allowing a single dollar-based commitment to cover usage globally. This reduces risk for smaller regional workloads and makes it easier to consolidate savings across distributed environments.

Key characteristics of Compute Savings Plans:

> Service-wide coverage: Applies to all compute usage across EC2, AWS Fargate, and AWS Lambda.
> Cross-platform flexibility: Allows you to switch between instance families, move workloads across regions, or adopt new compute types without losing coverage.

Automatic discount application: Discounts are applied to any eligible compute usage up to the committed spend amount, simplifying management across multi-account or multi-service environments.
Term and payment options: Available in one- or three-year terms with All Upfront, Partial Upfront, or No Upfront payment models.
Best for: Ease of management, organizations with dynamic infrastructure yet consistent spend, containerized or serverless architectures, or evolving compute strategies that require horizontal flexibility without sacrificing savings.

Amazon SageMaker Savings Plans

Amazon SageMaker Savings Plans are a separate type of AWS Savings Plan designed specifically for Amazon SageMaker usage. They offer up to 64% savings compared to On-Demand rates and apply to a broad set of SageMaker components, including Studio-Notebook, Training, BatchTransform, etc.

These plans work similarly to other Savings Plans in that they are based on a fixed hourly spend and apply automatically to eligible usage. However, they are limited to SageMaker services and do not cover general-purpose EC2, Fargate, or Lambda compute.

Key characteristics of SageMaker Savings Plans:

SageMaker-specific scope: Applies only to SageMaker services such as Studio Notebooks, Training Jobs, Batch Transform, Hosting Services, and Processing Jobs.
Automatic application: Discounts are applied to any eligible SageMaker usage in real time, up to the committed hourly spend.
Term and payment options: Available in one- or three-year terms with All Upfront, Partial Upfront, or No Upfront options.

Best for: Teams running consistent machine learning workloads on SageMaker with predictable usage across training and inference.

How To Purchase AWS Savings Plan?

AWS have recently released a helpful tool called Purchase Analyzer to gauge the impact of a Savings Plan Purchase. Users may set a lookback period for which Purchase Analyzer will calculate the cost, coverage, and utilization of a planned Savings Plan purchase. It is recommended to use this tool prior to committing to a SP purchase.

Here's a step-by-step overview to guide your AWS Savings Plan purchase through AWS console:

Sign in to the AWS Management Console: Navigate to the Billing and Cost Management section, then select "Savings Plans" from the left-hand menu. This will take you to the Savings Plans dashboard.
Choose the type of Savings Plan: Click "Purchase a Savings Plan" at the top; you'll be prompted to select between Compute Savings Plans, EC2 Instance Savings Plans, or SageMaker Savings Plans. Make this choice based on your flexibility needs and workload predictability.
Specify your commitment details: Enter your hourly spend commitment, select a one- or three-year term, and choose a payment option. For EC2 Instance Savings Plans, you'll also need to specify the instance family and region. Optionally, you may also specify a start date and time for coverage to take effect.
Confirm and purchase: After reviewing your configuration, Add to cart and finalize the purchase by clicking "Submit order." Once active, the Savings Plan will automatically apply to eligible usage across your environment.

Returning a Newly Purchased AWS Savings Plan

AWS has now introduced a seven-day return window for newly purchased Savings Plans, allowing customers to reassess their commitments without being locked into a long-term decision.

Savings Plans are eligible for return if they have an hourly commitment of $100 or less and were purchased within the last seven days. Additionally, the return must be processed within the same calendar month of purchase, according to UTC time.

Upon returning a Savings Plan, users receive a 100% refund for any upfront payments made towards the plan, and it will be reflected in the customer's bill within 24 hours of the return. Any usage that was previously covered by the plan will either be charged at AWS's standard On-Demand rates or covered by another applicable Savings Plan, if available.

However, there are some restrictions to returning Savings Plans you should be aware of. You can learn more about the restrictions of the 7-day return window here.

AWS Savings Plans vs Reserved Instances (RIs)

Both AWS Savings Plans and Reserved Instances offer substantial savings compared to On-Demand pricing, but they differ in flexibility, scope, and how commitments are applied. While RIs require you to commit to specific instance attributes like type, region, and OS, Savings Plans use a dollar-per-hour commitment model, which, in the case of Compute Savings Plans, applies more broadly across services and configurations.

Understanding the differences between these instruments is essential for choosing the right one based on your workload predictability, infrastructure strategy, and operational risk tolerance. The table below breaks down how each option compares across key dimensions.

Features	Most restrictive; highest discount; possible to transfer/sell	Simple; highest discount	Most control but requires effort	Simple; most flexible
Commitment Unit	Instance	Dollar	Instance	Dollar
Avg Discount (1 year)	38%	29%	38%	29%
Avg Discount (3 years)	58%	58%	51%	51%
Service	EC2 only	EC2 only	EC2 only	EC2 + Fargate + Lambda
Geography	Region specific	Region specific	Region specific	Spans all regions
Instance Family	Fixed	Fixed	Fixed	Automatically flexes

Instance Size	Automatically flexes for Linux only; others fixed	Automatically flexes	Automatically flexes for Linux only; others fixed	Automatically flexes
Operating System	Fixed	Automatically flexes	Fixed	Automatically flexes
Tenancy	Fixed	Automatically flexes	Fixed	Automatically flexes
Term	1 or 3 years; shorter terms from RI Marketplace	1 or 3 years only	1 or 3 years; shorter terms via exchanges	1 or 3 years only
Payment options	No, Partial, All Upfront	No, Partial, All Upfront	No, Partial, All Upfront	No, Partial, All Upfront
Recourse if unutilized	Attempt to sell on RI Marketplace	None	Exchange for another instance size, family, operating system, tenancy, term, or payment option	None

For a more detailed overview, check out our blog: AWS Savings Plans vs Reserved Instances.

Best Practices To Make the Most of AWS Savings Plans

Maximizing the impact of AWS Savings Plans requires more than just a well-timed purchase. It involves data-driven forecasting, continuous monitoring, and the flexibility to adjust as your infrastructure evolves. Below

are 10 best practices to help you extract the full value of your SP investment.

Use historical spend to size your commitment

Start by analyzing three to six months of your spend through AWS Cost Explorer, the Cost and Usage Report (CUR), or Savings Plan Purchase Analyzer. Filter out spiky workloads and focus on consistently running baselines. These are ideal for anchoring your hourly commitment without exposing yourself to unnecessary risk.

Choose the right Savings Plan type for your needs

If your environment is largely EC2-based and stable within a specific family and region, EC2 Instance Savings Plans offer better discounts. For teams with containerized, serverless, or cross-region workloads, Compute Savings Plans provide flexibility without compromising on automation or savings.

Don’t overcommit upfront, start with core usage

While the temptation to maximize discounts is high, overcommitting leads to wasted spend. Focus your initial SP purchases on the portion of usage you’re confident won’t decline. As coverage data improves over time, you can incrementally increase your commitment.

Adopt a rolling commitment strategy

AWS advises using a rolling strategy for AWS SPs, which says that instead of purchasing a large Savings Plan in one transaction, spread your commitment across multiple purchases over several months. This staggers expiration dates and creates natural checkpoints to evaluate whether to renew, expand, or reduce commitments based on real usage.

Monitor coverage and utilization continuously

Although SPs apply automatically, they are not a “set and forget” tool. Monitor whether usage is consistently covered or if you’re paying On-Demand for workloads that could be discounted. Look out for spikes above your committed amount or drops that leave commitment value unused.

Understand what’s covered and what’s not

Savings Plans do not apply to services like RDS, ElastiCache, or Redshift. If your environment spans multiple service categories, make sure you’re not assuming full coverage. This is especially important when reporting on savings performance across business units.

Align payment options with financial strategy

All Upfront provides the deepest discount but requires a large capital outlay. If budget flexibility is more important, consider Partial or No Upfront options. Match your payment model to your organization’s cash flow and budgeting cycle, not just the raw discount.

Track Effective Savings Rate (ESR) and Commitment Lock-in Risk (CLR)

Utilization and coverage are surface-level metrics, they show usage and coverage efficiency but not true financial outcomes. Effective Savings Rate (ESR) tells you how much you’re actually saving, while Commitment Lock-in Risk (CLR) indicates the duration, in months, of vendor lock-in due to long-term commitments. Monitoring both provides a more complete view than utilization metrics alone, helping FinOps teams assess whether they’re optimizing safely.

Re-evaluate commitment levels before renewal

Usage patterns change. Before renewing expiring SPs, review how your compute footprint has evolved. Factor in workload shifts, migrations (e.g., from EC2 to Fargate), or architectural changes to avoid repeating a commitment that no longer aligns with business reality.

Automate SP management to maintain optimization at scale

Savings plan management isn't a one-time task. Usage patterns shift daily, new instance types are released, and project demands evolve, requiring constant evaluation and expert decision-making. But managing this manually isn't scalable. Your usage is dynamic; your commitments are not.

To optimize SPs without distracting engineering and FinOps teams, businesses turn to autonomous platforms like ProsperOps, which dynamically adjust commitments to maximize savings while minimizing lock-in risk, all hands-free.

Recently, our Senior Product Manager, Clay Wolcott, co-hosted a webinar with Adam Richter, Senior Optimization Solutions Architect at AWS. Together, they explored advanced strategies for leveraging AWS Savings Plans to reduce costs and improve flexibility. Adam also shared several expert insights on commitment management.

Catch the full webinar and key takeaways here.

Simplify AWS Savings Plan Management With ProsperOps

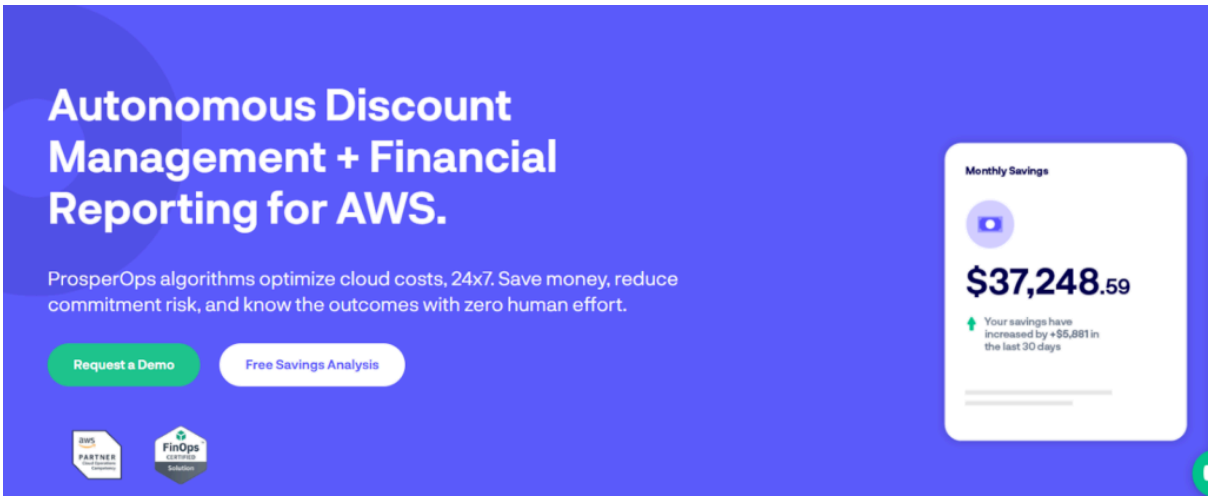

Managing Savings Plans manually can be time-consuming, complex, and risky, especially as your usage patterns evolve. To maximize savings without constant oversight, organizations are turning to automation. That's where ProsperOps comes in.

ProsperOps automates AWS Reserved Instance and Savings Plan management and dynamically optimizes your commitments in near real time to ensure maximum savings with minimal risk. A key part of this strategy is Savings Plan Adaptive Laddering, which distributes commitments over time instead of in large batches. This creates regular expiration points, reducing lock-in risk and giving your team the flexibility to respond to changing usage.

We help businesses build a flexible commitment portfolio that adapts to dynamic usage patterns. With near real-time data ingestion, ProsperOps ensures maximum savings without manual intervention, complex forecasting, or operational overhead.

This hands-free approach to cloud cost optimization can save your team valuable time while ensuring automation continually optimizes your AWS, Azure, and Google cloud discounts for maximum Effective Savings Rate (ESR).

In addition to autonomous rate optimization, ProsperOps now supports usage optimization through its resource scheduling feature, ProsperOps Scheduler. Our customers using Autonomous Discount Management™ (ADM) can automate resource state changes on weekly schedules to reduce waste and lower cloud spend.